

Docker: an insider view

Google Developers Group Meetup
November 2013, Google West Campus 2



Michael Crosby — @crosbymichael

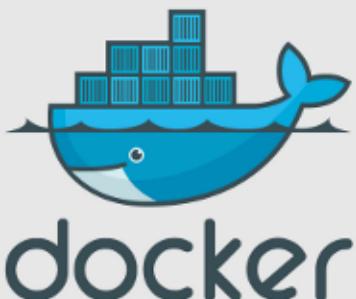
Jérôme Petazzoni — @jpetazzo

Victor Vieux — @vieux

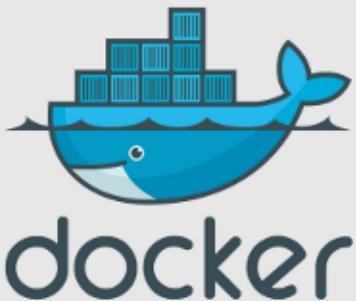


Outline

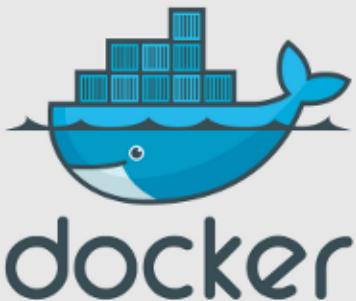
- what is Docker
- why it is written in Go
- some implementation details
- drawbacks
- more drawbacks
- discussion



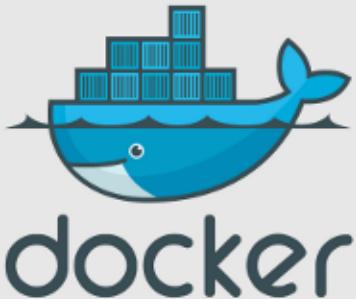
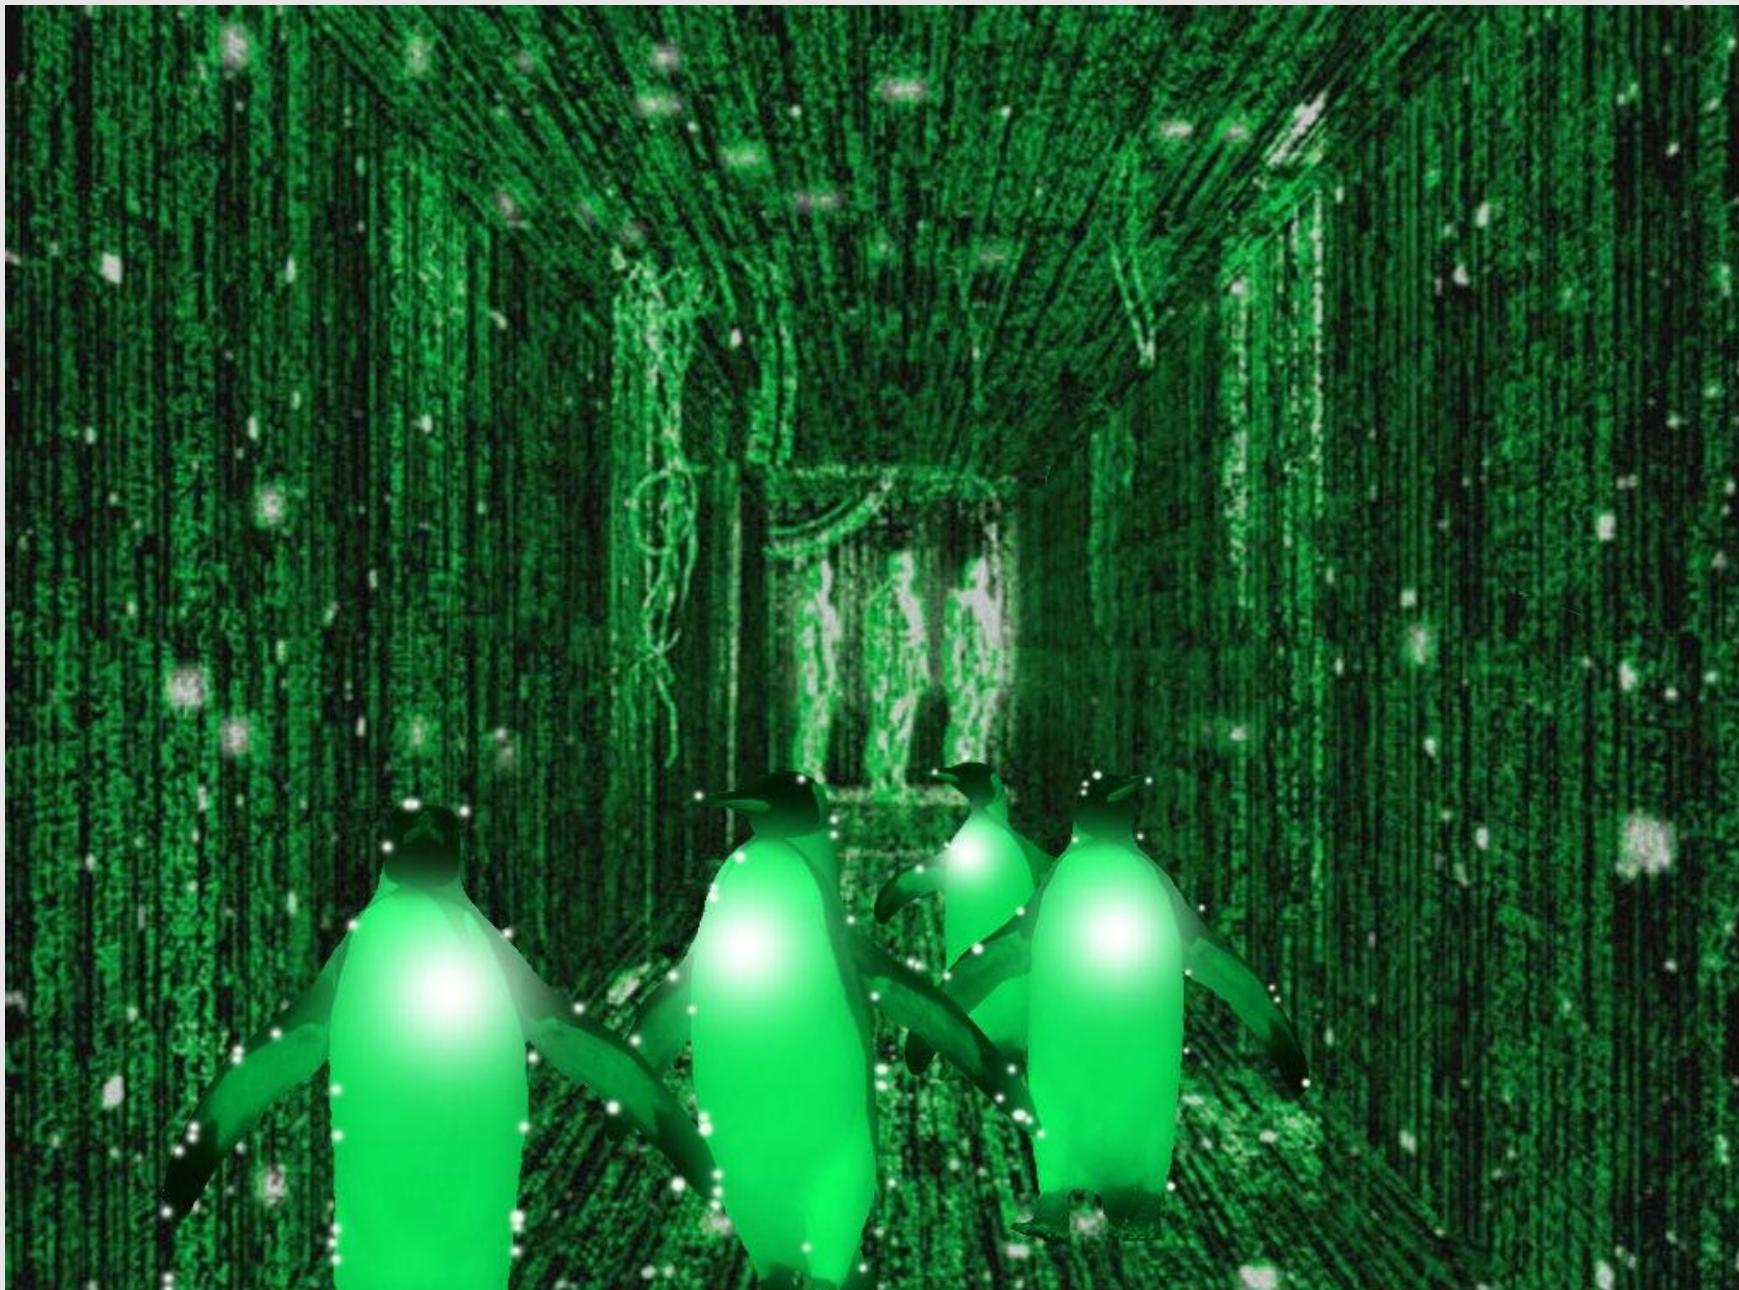
What's Docker?



Docker solves the “Matrix from Hell”

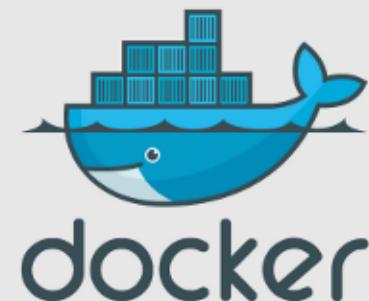


Docker solves the “Matrix from Hell”

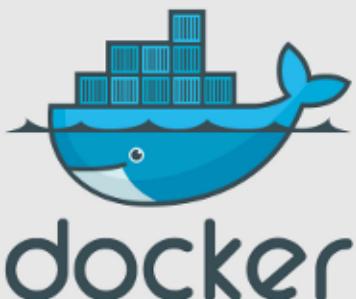


Docker solves the “Matrix from Hell”

static website	?	?	?	?	?	?
web frontend	?	?	?	?	?	?
background workers	?	?	?	?	?	?
user DB	?	?	?	?	?	?
analytics DB	?	?	?	?	?	?
queue	?	?	?	?	?	?
	dev VM	QA server	single prod server	on-site cluster	public cloud	contributor laptop

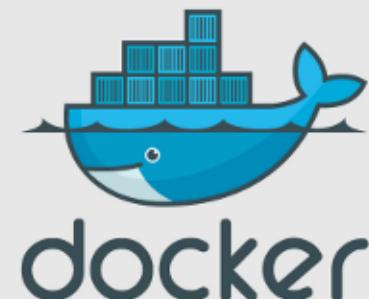


Real-world analogy: containers



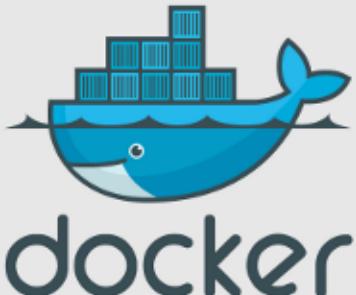
Worldwide shipping: another Matrix from Hell

clothes	?	?	?	?
wine	?	?	?	?
coffee	?	?	?	?
electronics	?	?	?	?
cars	?	?	?	?
raw materials	?	?	?	?
	ships	trains	trucks	storage

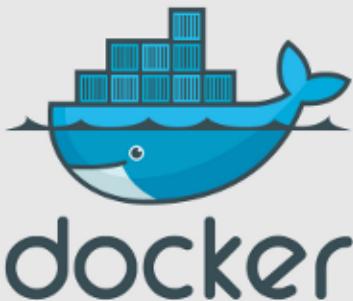


Solution: the *intermodal shipping container*

clothes				
wine				
coffee				
electronics				
cars				
raw materials				
	ships	trains	trucks	storage



Solution to the deployment problem: the Linux container



What's a Linux container?

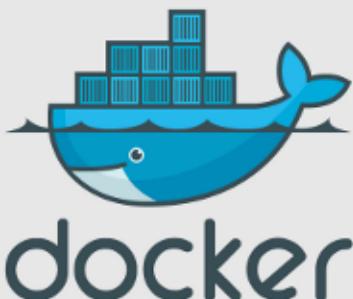
High level approach

Chroot on steroids

- normal processes, but isolated
- share kernel with the host
- no device emulation (neither PV nor HVM)
- doesn't need a /sbin/init
(runs the app/DB/whatever directly)



“Application Container”



What's a Linux container?

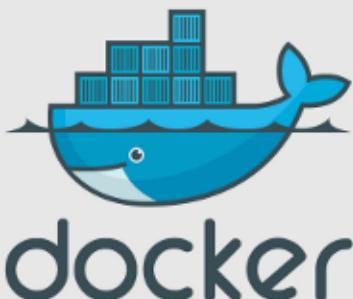
Low level approach

Lightweight Virtual Machine

- own process space
- own network interface
- can run stuff as root
- can have its own /sbin/init
(different from the host)



“Machine Container”

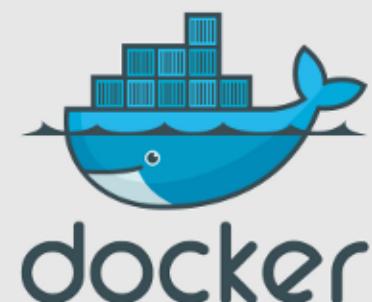


Separation of concerns: Dave, from the dev team

- my code
- my framework
- my libraries
- my system dependencies
- my packaging system
- my distro
- my data



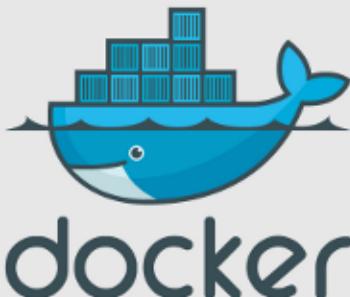
I don't care where it's running or how.



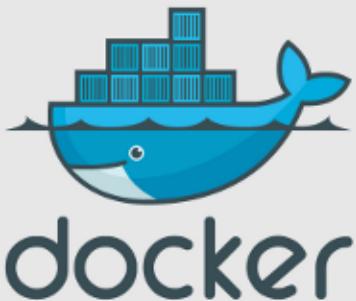
Separation of concerns: Oscar, from the ops team

- logs
- backups
- remote access
- monitoring
- uptime

I don't care what's running in it.



OK, so what's Docker?



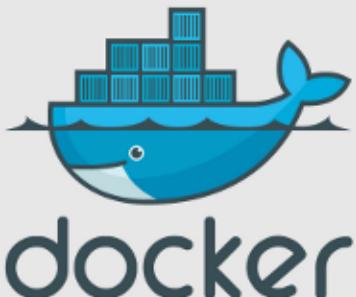
Runtime for Linux containers

```
jpetazzo@tarrasque:~$ sudo docker run -t -i ubuntu bash  
root@092ee318746f:/#
```

```
jpetazzo@tarrasque:~$ sudo docker run -d john costa/redis  
c6000fa5ddc6
```

```
jpetazzo@tarrasque:~$ sudo docker inspect c6000fa5ddc6  
...
```

```
  "NetworkSettings": {  
    "IPAddress": "172.17.0.8",  
    "Ports": {  
      "6379/tcp": null  
    }  
  }  
...
```



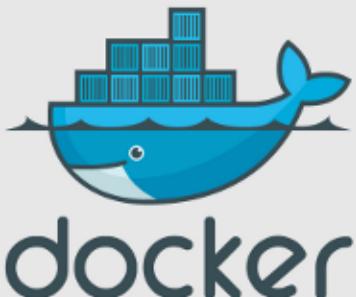
Standard format for containers, and a place to share them

- fetch an image from the registry with “docker pull”
- enter the image with “docker run”, do some changes
- record those changes with “docker commit”, repeat as many time as needed
- then share the result with “docker push”, on the public registry or on a private one



Why Go?

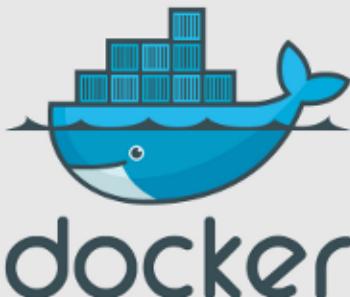
The Five Reasons Why We Used Go



Why Go?

1) static compilation

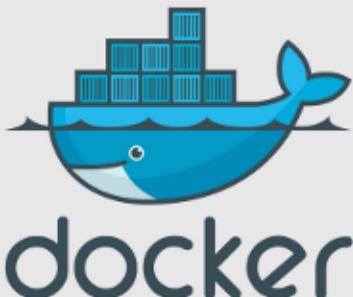
- “go build” will embed everything you need
(no more “install this in order to run my stuff”)
- ... except dynamic libraries if you use cgo
(cgo lets you use any C library)
- and ... except libc
(but who doesn’t have libc?)
- you can have a real static binary
(if you hack the build process a bit...)



Why Go?

1) static compilation

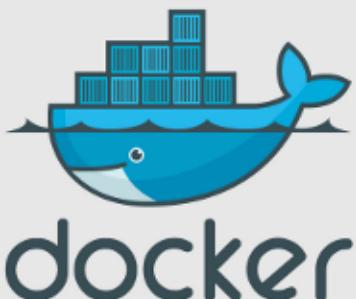
- easier to install, easier to test, easier to adopt
- good candidate for bootstrap
(i.e. the first installed thing,
which will install the other things)
- “no dependencies? me gusta!”
-- *Anonymous BOFH*



Why Go?

2) neutral

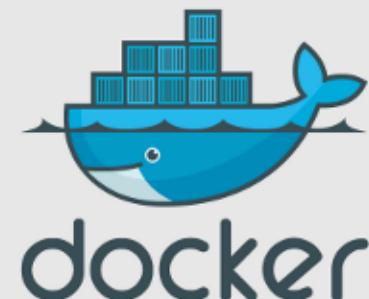
- it's not C++
- it's not Python
- it's not Ruby
- it's not Java



Why Go?

3) it has what we need

- good asynchronous primitives
(wait for I/O, wait for processes...)
- low-level interfaces
(manage processes, syscalls...)
- extensive standard library and data types
- strong duck typing

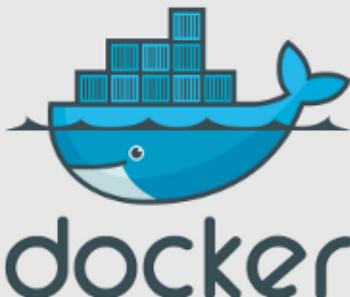


Why Go?

4) full development environment

Go addresses multiple issues of the development workflow.

- go doc (see documentation for any package)
- go get (fetch dependencies on github etc.)
- go fmt (solve “tabs vs. spaces” once for all)
- go test (runs all `Test*` functions in `*_test.go`)
- go run (rapid script-like prototyping)



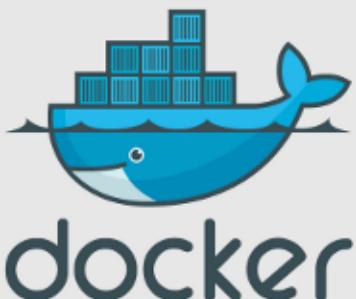
Why Go?

5) multi-arch build

...without pre-processors

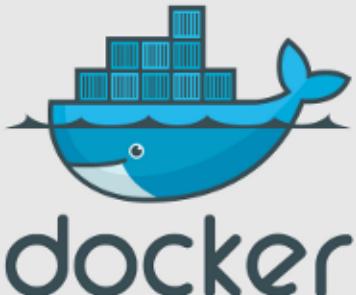
`_linux.go`

`_darwin.go`



Drawbacks

Because that's the
most interesting thing
in Amazon reviews



“Go doesn’t solve any problem”



[Jonn Mostovoy](#) @podmostom 29 Sep

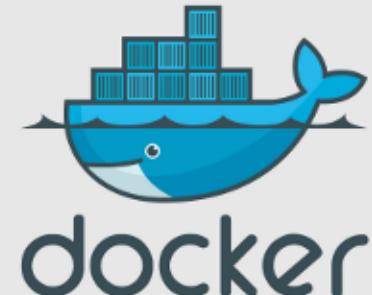
@benoitc @arnaudsj in not solving any problems.

Everything they claim to solve is better solved in Dv2.

0, Erlang or will be solved in Rust..

But...

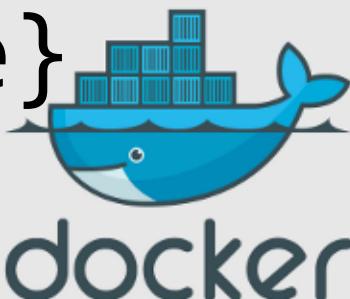
- Go is easier than Erlang
- Go is more real than Rust
- Don't know about D 2.0 though!



maps aren't thread-safe

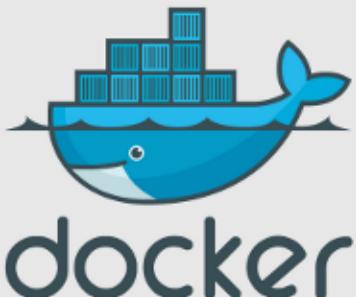
- deliberate decision: they're *fast* (and it's up to you to ensure that they're *safe*)
- you have to protect access with `sync.Mutex`
- or, use channels of channels!

```
m := NewMap()  
response := make(chan string)  
m<-Get{Key: "fortytwo", Replyto: response}  
value := <-response
```



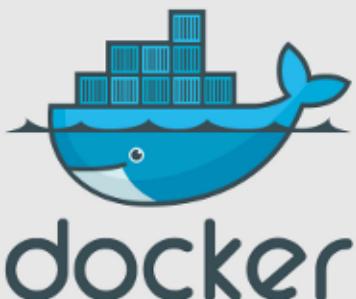
go get

- can't pin a particular revision
- Docker had to vendor all dependencies
(i.e. import their source code in our repo)
- must deal with private repos manually



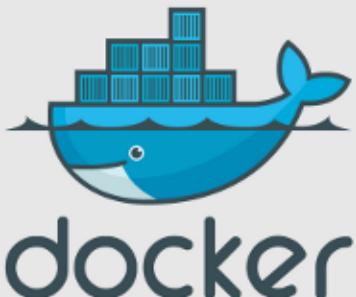
go test

- can't have destructors/cleanups/... in tests
- use a test named “z_final_test.go”
- ... which doesn't work too well when running individual tests!



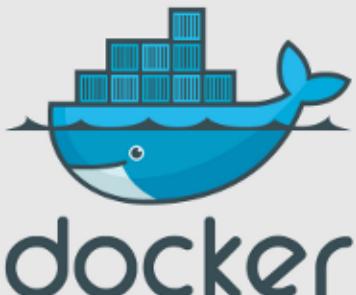
go build

- it's painful to build multiple binaries, when they share some common code
- each program has to be in its own package, and package must be “main”
- you *have to* put shared/common stuff aside, or use named import tricks (bleh)



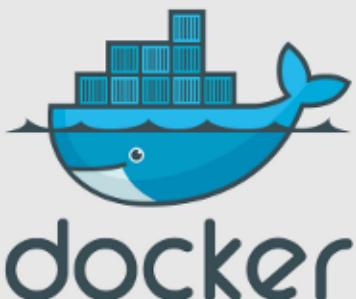
flag package

- doesn't handle short/long options
(-o --option)
- doesn't handle options grouping
(-abc -a -b -c)
- seriously just don't use it
(use getopt or go-flags instead)
unless your primary target is Plan9 of course!



no IDE

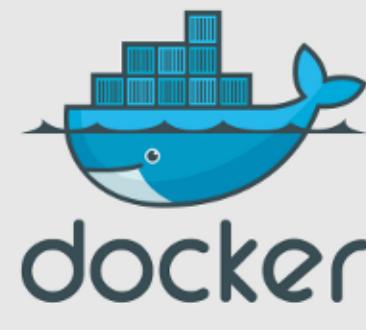
Let's ask Samuel Jackson
what he thinks about it



EMIACS



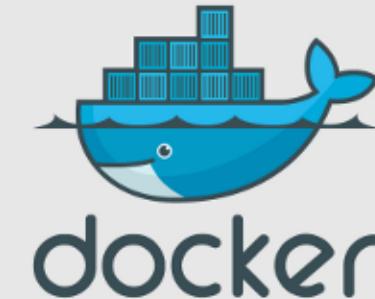
memegenerator.net



**AM I THE ONLY ONE AROUND
HERE**

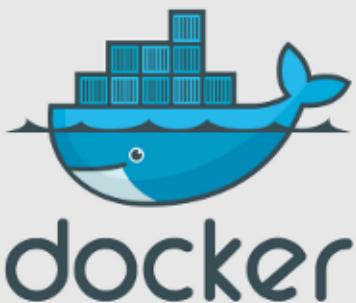
TO USE WIMP

memegenerator.net



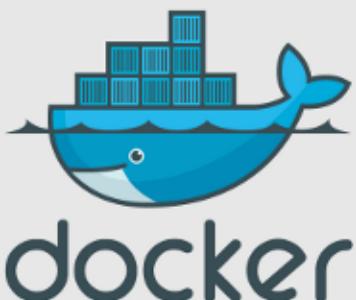
Error handling can be verbose

```
if err := openFile(); err != nil {  
    return err  
}  
if err := readAll(); err != nil {  
    return err  
}  
if err := closeFile(); err != nil {  
    return err  
}  
...
```



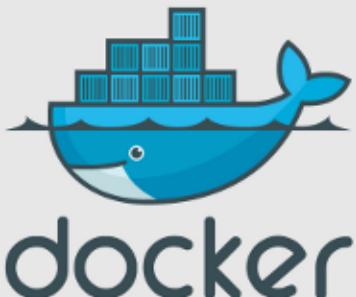
Error handling

- panic / recover
(don't abuse it; only use it internally!)
- deploy one-off utility types
<http://talks.golang.org/2013/bestpractices.slide#5>



can't select on readers/writers

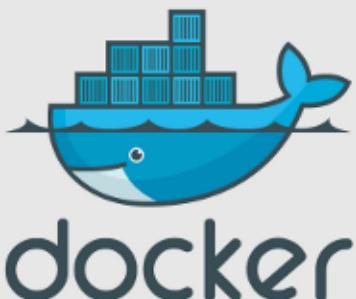
```
select {  
    case data = <-inputChannel:  
        ...  
    case data = <-otherInputChannel:  
        ...  
    case data = connection.Read():  
        ...  
}
```



can't select on readers/writers

- for readers:
 - start a goroutine with access to an output channel
 - the goroutine do blocking reads on the readers...
 - ...and pushes the results on the output channel
- for writers:
 - adaptation is left as an exercise for the reader

→ reduce everything to channels!



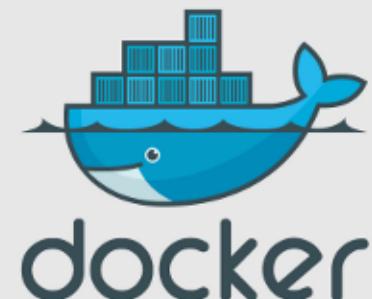
Some useful resources:

- The Holy Reference:
<http://golang.org/ref/spec>
- Some useful patterns:
http://golang.org/doc/effective_go.html
- Some best practices:
<http://talks.golang.org/2013/bestpractices.slide>
- Embedded documentation:
`godoc -http=:6060 & chrome http://localhost:6060`

Michael Crosby — [@crosbymichael](https://twitter.com/crosbymichael)

Jérôme Petazzoni — [@jpetazzo](https://twitter.com/jpetazzo)

Victor Vieux — [@vieux](https://twitter.com/vieux)



Thank you! Questions?

Michael Crosby — @crosbymichael

Jérôme Petazzoni — @jpetazzo

Victor Vieux — @vieux

